



SIGGRAPH THINK
BEYOND
2020 [S2020.SIGGRAPH.ORG](https://s2020.siggraph.org)



JUNE 2020 WEBINAR

Hands-on Workshop: Machine Learning and
Neural Networks

SIGGRAPH NOW



RAJESH SHARMA

SOFTWARE ENGINEER

Walt Disney Animation Studios



Machine Learning

————— Rajesh Sharma —————

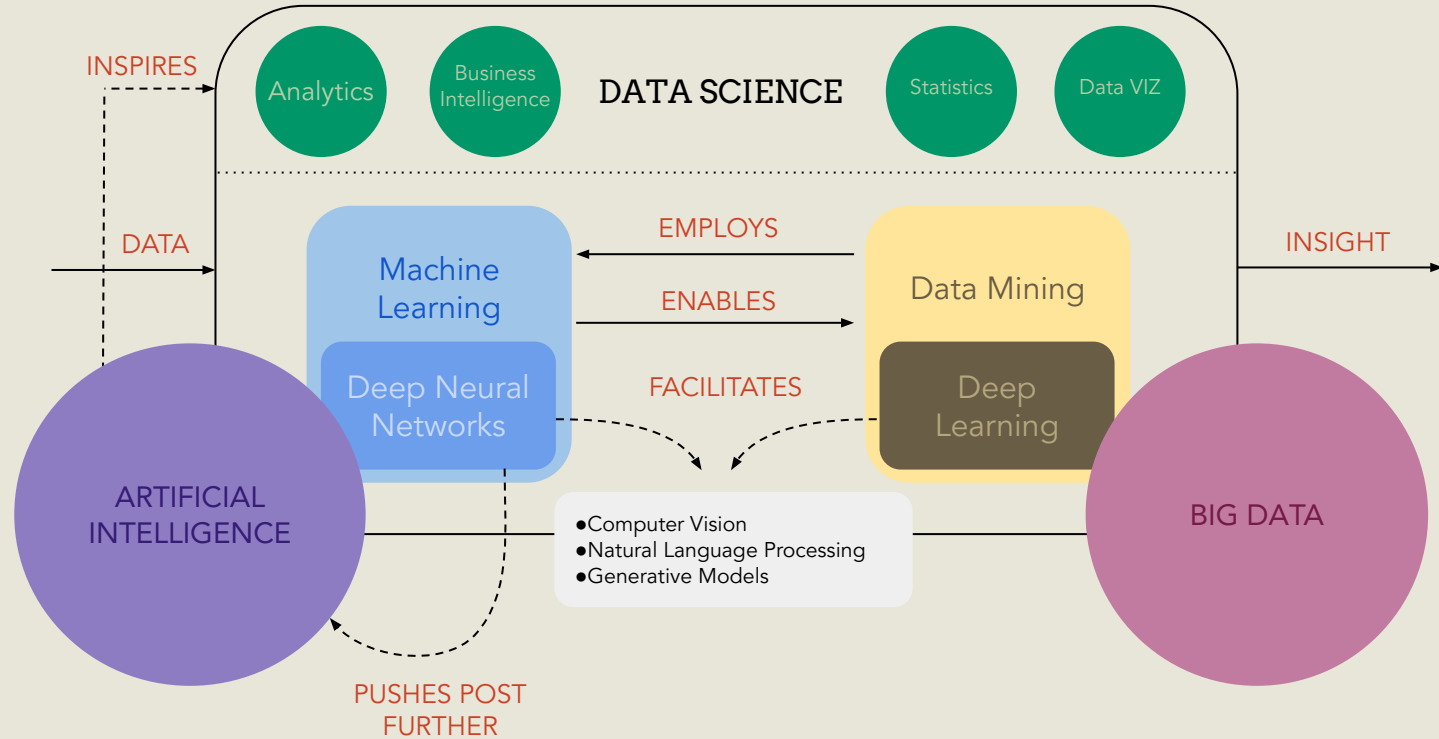
Today

- Introduction
- What is Machine Learning?
- What are Neural Networks?
- How do I 'solve' a neural network?
- What are Autoencoders?
- What is CNN?

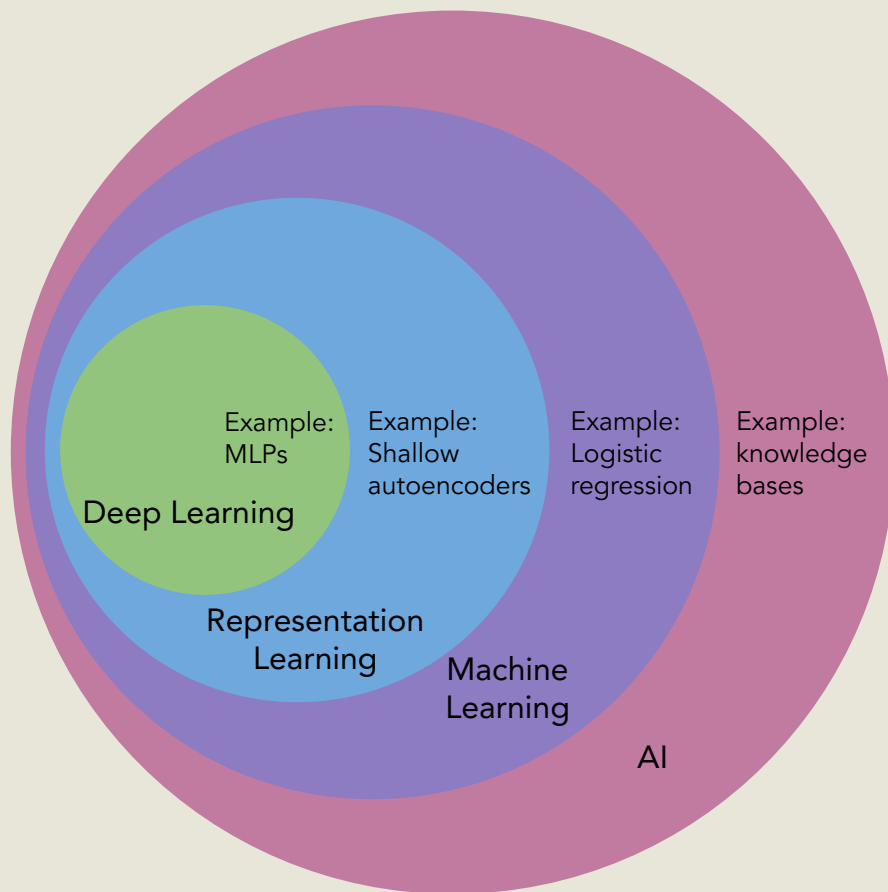
Housekeeping

- Link to today's slides and Colab notebooks:
 - Log in to your google drive
 - Find the shared folder 'Disney Machine Learning Webinar'
- Use the 'Ask a question' feature
- After the lecture: @xarmalarma, #siggraphNOW

The BIG Picture



AI & Machine Learning



What problem are you solving?

Question	AI/ML Task	Healthcare	Retail	Finance
Yes or No?	Detection	Cancer Detection	Targeted Ads	Cybersecurity
What type ?	<i>Classification</i>	Image Classification	Basket Analysis	Credit Scoring
What size ?	Segmentation	Tumor size	Customer Types	Risk Analysis
What result ?	<i>Prediction</i>	Survivability	Sentiment/Behavior	Fraud Detection
What action ?	Recommendation	Therapy	Recommendation	Fast Trading


Not a solution to every type of problem

When NOT to use	When to use
<p>✗ Can use computations, or algorithms or simple rules that can be programmed</p>	<p>✓ Problem cannot be solved using rule-based solutions</p>
	<p>✓ Model is complex or has too many factors</p>
	<p>✓ Need to scale to large number of inputs or factors</p>


How

- FRAMING: What is observed & what answer you want to predict
- DATA COLLECTION: Collect, clean, and prepare data
- DATA ANALYSIS: Visualize & analyze the data
- FEATURE PROCESSING: Transform raw data for better predictive input
- MODEL BUILDING: Design and build the learning algorithm
- TRAINING: Feed data to the model and evaluate the quality of the models
- PREDICTION: Use model to generate predictions for new data instances


Example (Linear Regression)




Example (Linear Regression)



Example (Regression)



Example (Regression) - Sum of least squares



Prediction: $\hat{y} = ax + b$

Actual: y_i

Error: $|y_i - \hat{y}_i|$

Total Squared Error:


$$\sum (y_i - \hat{y}_i)^2, \text{ for } i=(1, n)$$

Minimize Total Squared Error:


$$E(a, b) = \sum (y_i - ax_i - b)^2$$

(a, b) are the parameters (weights)

Regression - Minimize Error (Cost) via Gradient Descent




Linear function as a Network & a Matrix op




Example


degree 1






degree 2




degree 20




Adding non-linearity via an activation function



Network Node == Artificial Neuron



Adding complexity via a layer:



Universal Approximation Theorem:

In the **mathematical** theory of **artificial neural networks**, the **universal approximation theorem** states^[1] that a **feed-forward** network with a single hidden layer containing a finite number of **neurons** can approximate arbitrary well real-valued **continuous functions** on **compact subsets** of \mathbf{R}^n .

But, No Free Lunch Theorem:

For optimization problems... if an algorithm performs well on a certain class of problems then it necessarily pays for that with degraded performance on the set of all remaining problems.

Solving the network

- *Set the initial weights of the network randomly*
- *Make a forward pass through the network and compute output*
- *Compare the output with expected result and compute loss*
- *Change the weights by a small amount (Gradient descent via back prop)*
- *Repeat until desired minimization of error (cost) is achieved*

Summarizing

- Given: Features (X, Attributes), Output (Y, Labels, Ground Truth): $Y=f(X)$
- Network (Model)
- Loss Function (Metric, Cost)
- Activation Function (adds non-linearity, Ex: sigmoid, ReLU)
- Training (fit, Optimization to minimize Loss Function)
- Evaluate (performance, correctness)
- Predict (Inference, on new data)

Computer Graphics Applications

★ *Scheduling Optimization*

★ *Character AI*

★ *Style Transfer*

★ *Slow Motion*

★ *Up-Res*

★ *Denoising*

★ *Story Sentiment*

★ *Rough to Fine*

★ *Body Tracking*

★ *Image Generation*

Hands-on (software and environment)

- ★ *We'll be using a python virtual environment: Colab*
- ★ *Colab: Jupyter derived python IDE with tensorflow support*
- ★ Software and tools:
 - Python 3.x - programming
 - Tensorflow 2.1.0 - machine learning
 - Numpy - numerical mathematics, linear algebra
 - Pandas - data analysis
 - Matplotlib - plotting
 - Seaborn - advanced plotting

Hands-on

- ★ *Log in to your google drive*
- ★ Find the shared folder 'Disney Machine Learning Webinar'
- ★ Make a copy of:
 - Housing.ipynb,
 - *Let's take a look at the Housing.ipynb*
 - *Analysis of data and possible transformations*

Hands-on

```
# -*- coding: utf-8 -*-
"""Housing.ipynb
"""
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
#-----DATA READING
filename = 'https://download.mlcc.google.com/mledu-datasets/california_housing_train.csv'
# read comma separated values (csv) from the file
csv_data = pd.read_csv(filename, sep=',')

print(csv_data.shape) # matrix of data
print(csv_data.head()) # first five data points
print(csv_data.head().transpose())
print(csv_data.describe()) # simple statistics about the data

print(csv_data[['latitude', 'longitude']]) # pick two columns to print



# plot two variables, the third (total bedrooms) is the size of the dots
sns.relplot(x='latitude', y='longitude', size='total_bedrooms', alpha=0.5, palette='muted',
data=csv_data)

sns.pairplot(csv_data.head(1000)[['longitude', 'latitude', 'total_bedrooms', 'median_house_value',
'population']])

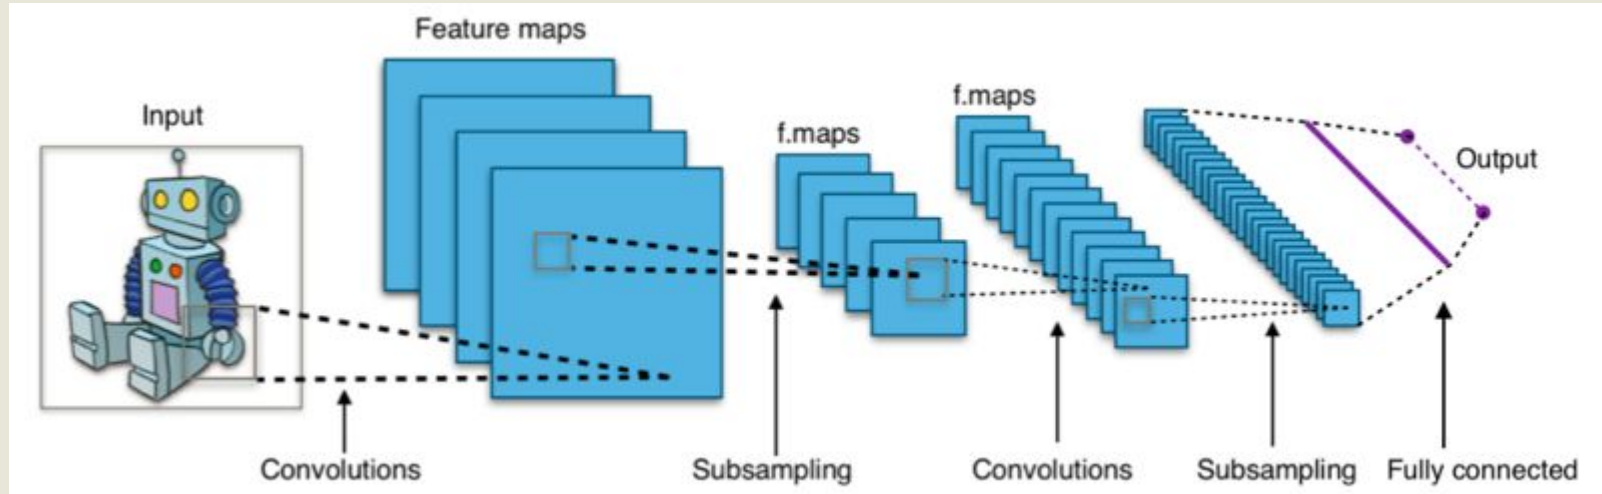
# pick random sample of data
sns.pairplot(csv_data.sample(n=1000)[['longitude', 'latitude', 'total_bedrooms',
'median_house_value', 'population']])
```

Kinds of Neural Networks


Auto Encoder (AE)



Kinds of Neural Networks (CNN)



Convolution (Extract High-Level Features)



Next Class

- Tensorflow
- Neural Net for Regression, Classification
- Homework:
 - Play with the data, try other plots
- @xarmalarma, #siggraphNOW

QUESTIONS?

Submit now!

THANK YOU

WANT TO HEAR MORE FROM SIGGRAPH AND DISNEY?

Check out SIGGRAPH Now on
YouTube



[YouTube.com/user/ACMSIGGRAPH](https://www.youtube.com/user/ACMSIGGRAPH)

Subscribe to the SIGGRAPH Spotlight
podcast

